

# CSE 291 – AI Agents Classical Control, Pre-deep Learning

Prithviraj Ammanabrolu

Thanks to David Silver's DeepMind RL Course and Sheila McIlraith's Planning Course at UofT. Some slides were adapted from there.

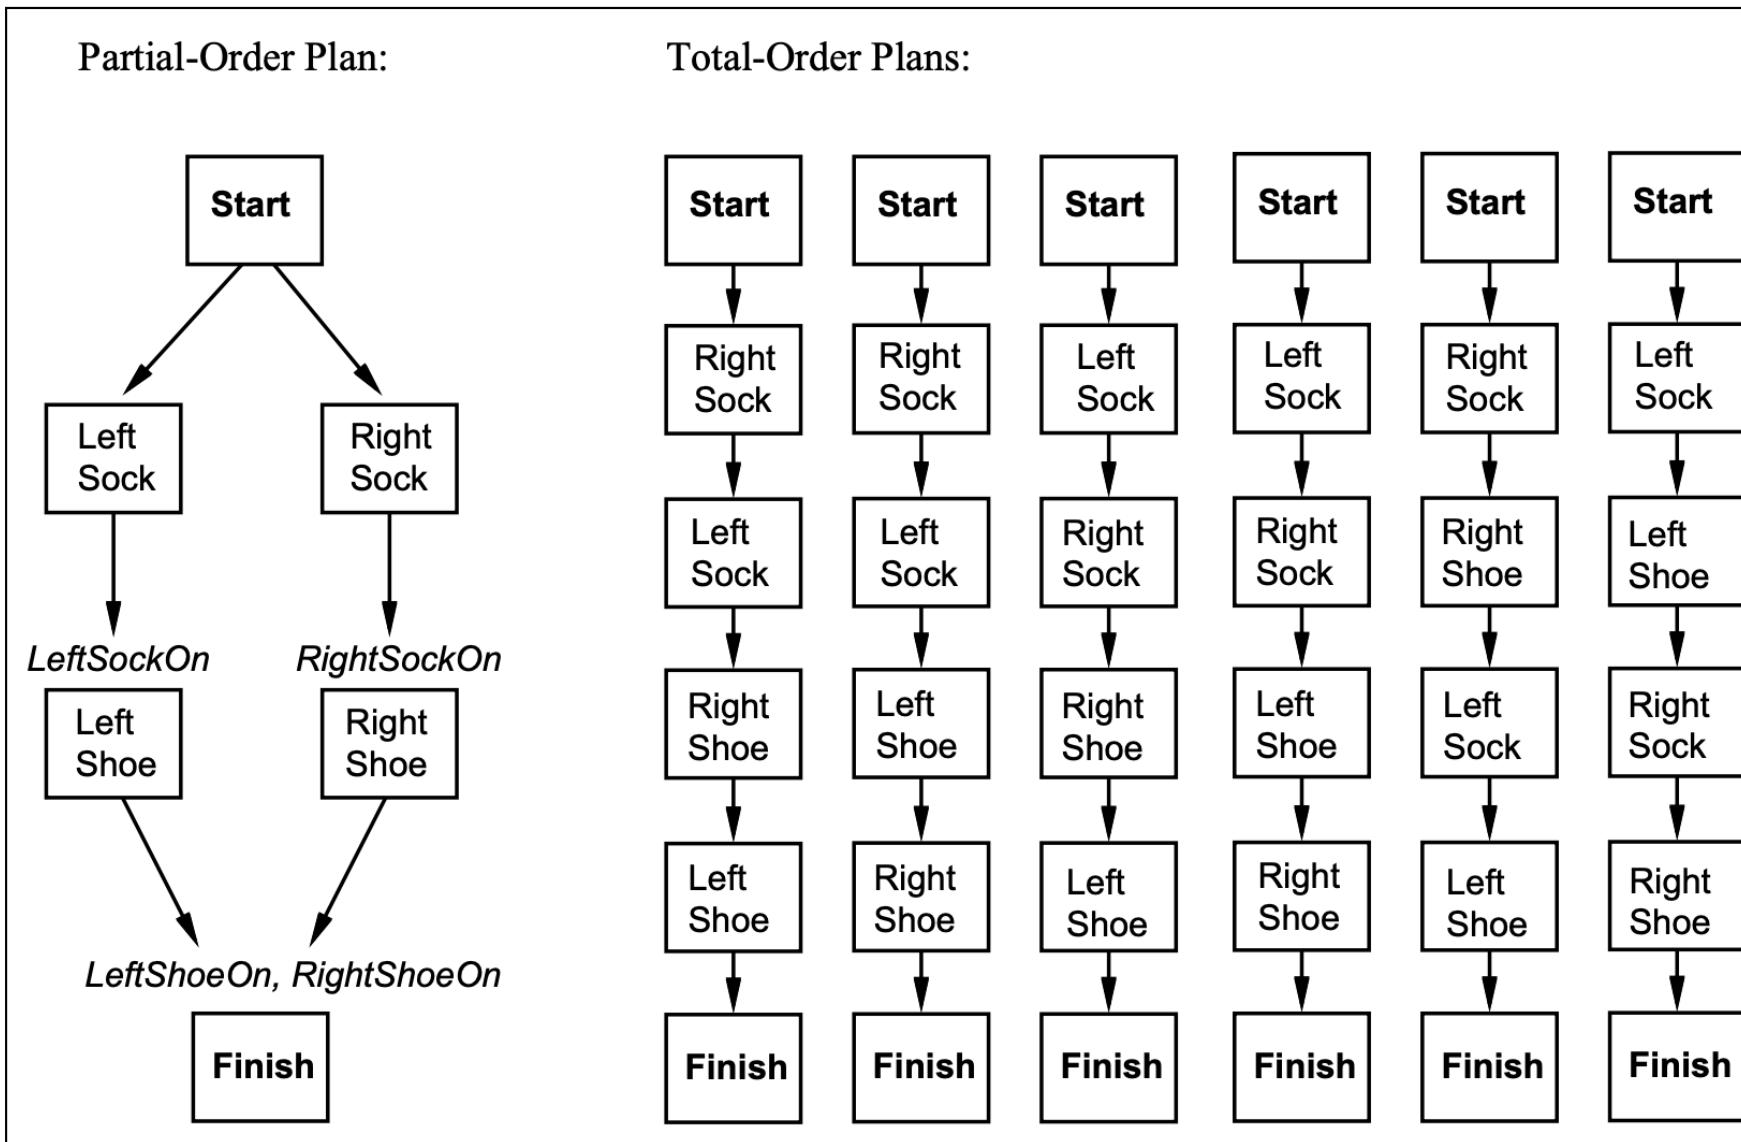
# Forward Search

- Some deterministic implementations of forward search:
  - breadth-first search
  - depth-first search
  - best-first search (e.g., A\*)
  - greedy search
- Breadth-first and best-first search are sound and complete But they usually aren't practical, requiring too much memory
  - Memory requirement is exponential in the length of the solution
- In practice, more likely to use depth-first search or greedy search
  - Worst-case memory requirement is linear in the length of the solution
  - In general, sound but not complete
  - But classical planning has only finitely many states
  - Thus, can make depth-first search complete by doing loop-checking

# Backward Search

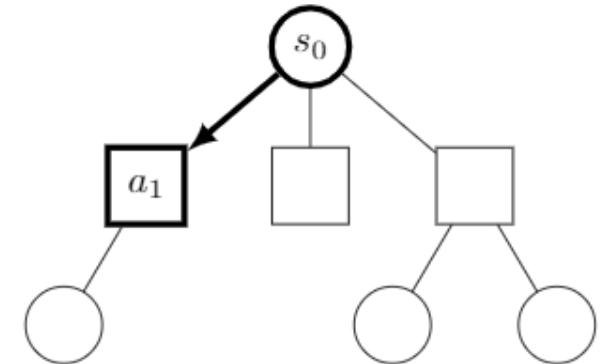
- For forward search, we started at the initial state and computed state transitions
  - new state =  $T(s,a)$
- For backward search, we start at the goal and compute inverse state transitions
  - new set of subgoals =  $T^{-1}(g,a)$
- To define  $T^{-1}(g,a)$ , must first define relevance: An action  $a$  is relevant for a goal  $g$  if
  - $a$  makes at least one of  $g$ 's literals true,  $g \cap \text{effects}(a) \neq \emptyset$
  - $a$  does not make any of  $g$ 's literals false,  $g + \cap \text{effects} - (a) = \emptyset$  and  $g - \cap \text{effects} + (a) = \emptyset$

# Total Order and Partial Order Plans



# Monte Carlo Tree Search

- 4 phases of building out and simulating paths along a search tree
- Various forms of this used in everything from Alpha Zero to modern LLM inference
- For arbitrary problem with start state  $s_0$  and actions  $a_i$
- All states have attributes:
  - Total simulation reward  $Q(s)$  and
  - Total no. of visits  $N(s)$



# Improvements to MCTS Components

- Improvements are possible for each of the parts I talked about
- Think about what it would take to improve selection / expansion phases

# Upper Confidence Trees (UCT)

- A way of improving the selection phase by treating selection as a multi-arm bandit problem: which possible action to select that maximizes the possible payout (reward) in the future

$$\text{UCT}(v_i, v) = \frac{Q(v_i)}{N(v_i)} + c \sqrt{\frac{\ln N(v)}{N(v_i)}}$$

# Upper Confidence Trees (UCT)

- A way of improving the selection phase by treating selection as a multi-arm bandit problem: which possible action to select that maximizes the possible payout (reward) in the future

$$\text{UCT}(v_i, v) = \boxed{\frac{Q(v_i)}{N(v_i)}} + c \sqrt{\frac{\ln N(v)}{N(v_i)}}$$

Exploit

# Upper Confidence Trees (UCT)

- A way of improving the selection phase by treating selection as a multi-arm bandit problem: which possible action to select that maximizes the possible payout (reward) in the future

$$\text{UCT}(v_i, v) = \frac{Q(v_i)}{N(v_i)} + c \sqrt{\frac{\ln N(v)}{N(v_i)}}$$

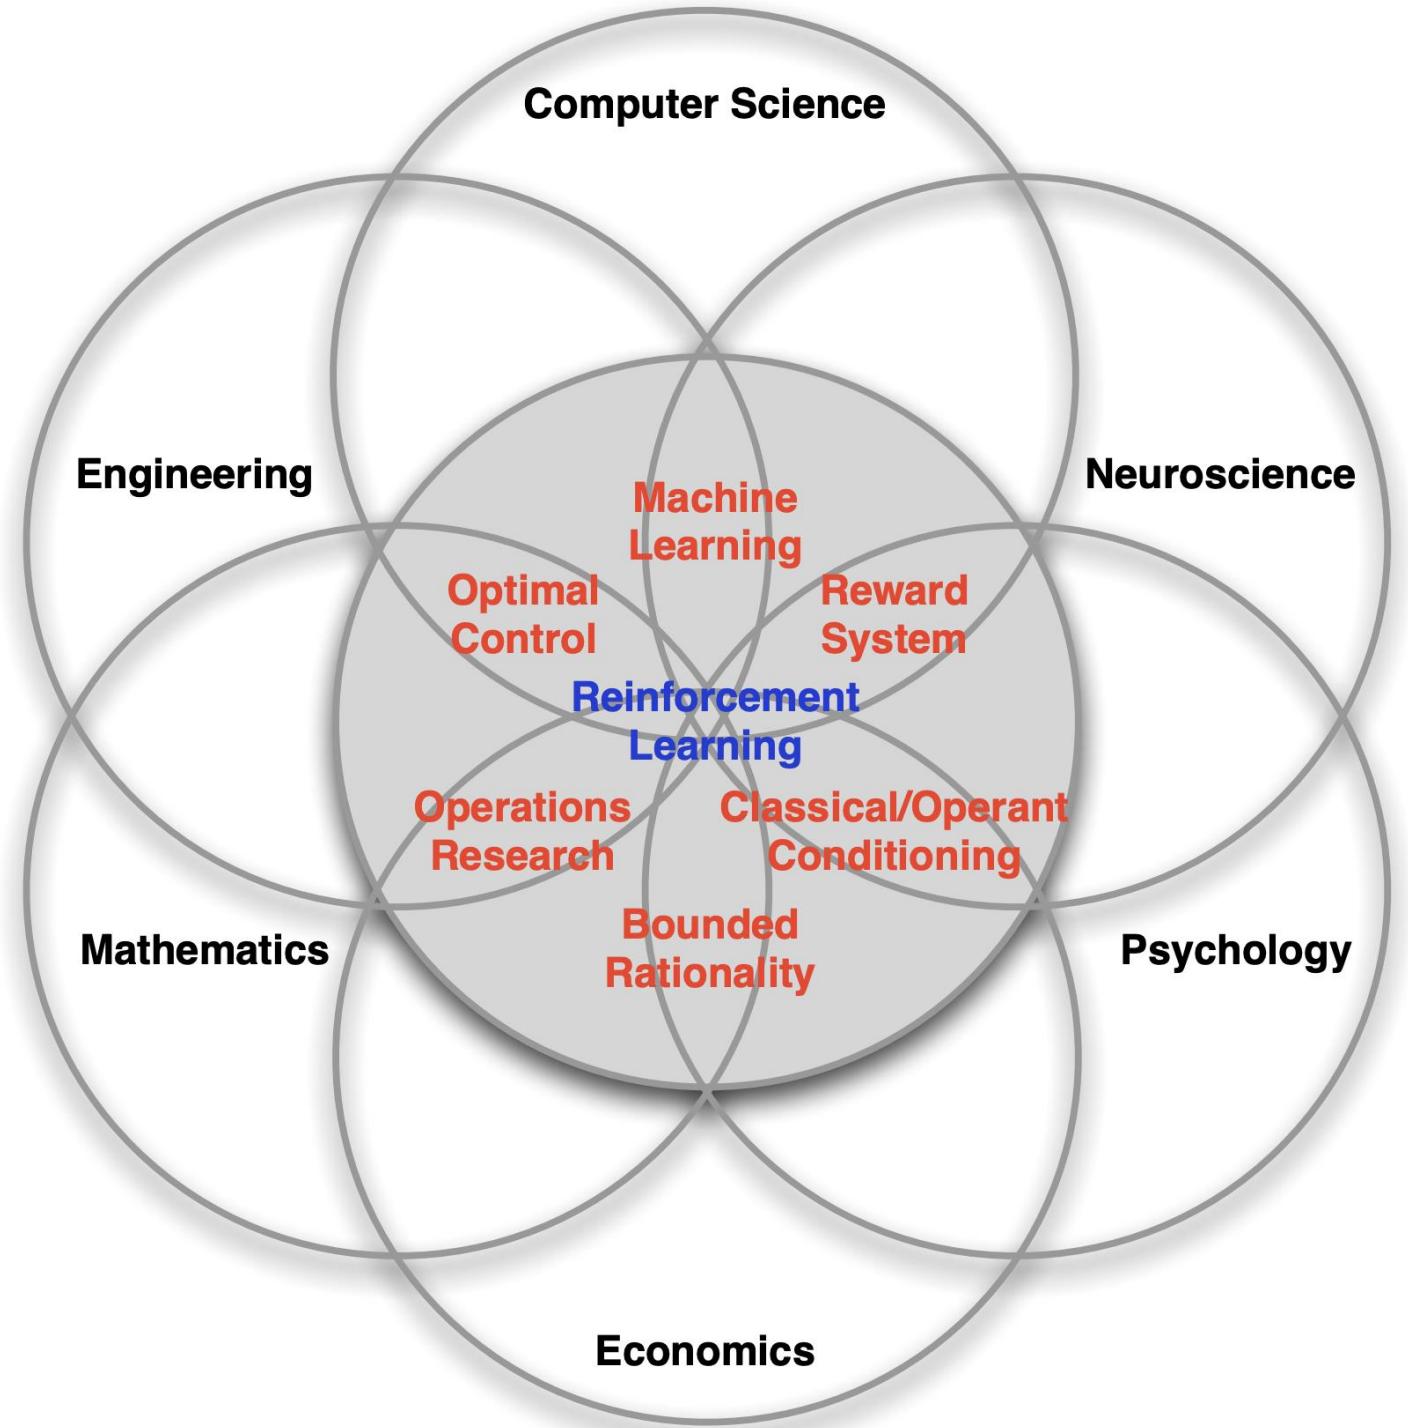
Exploit

Explore

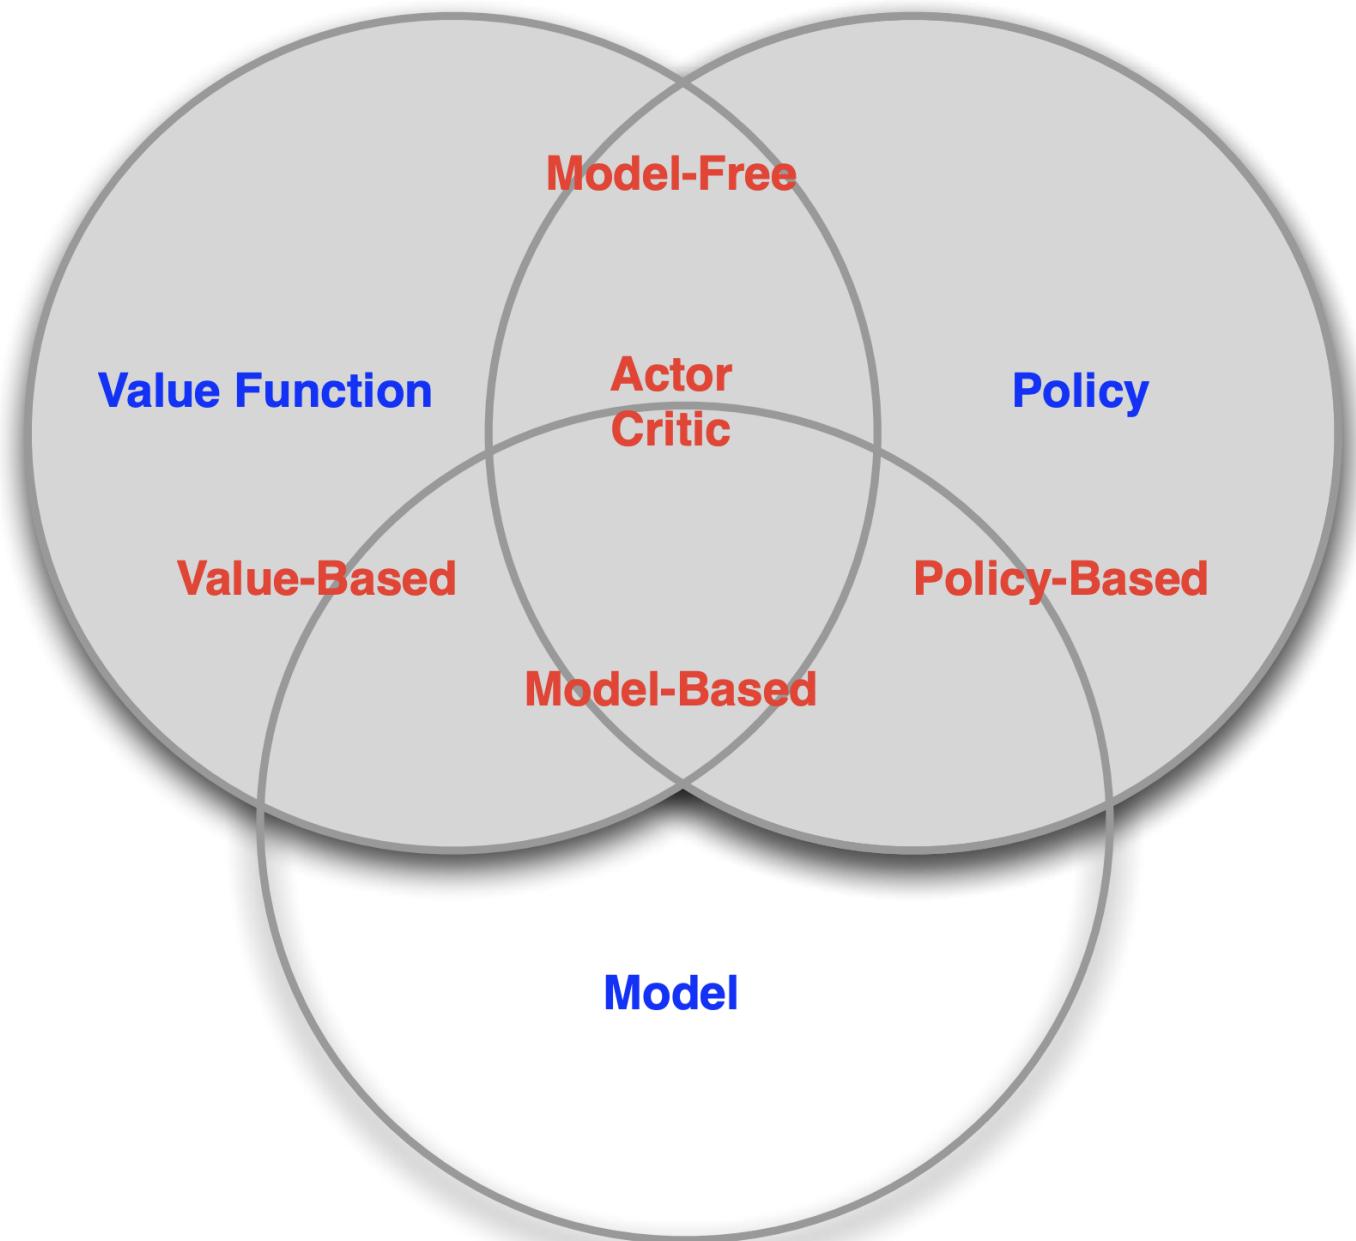
# Why Reinforcement Learning?

- Reinforcement Learning:
  - The environment is initially unknown
  - The agent interacts with the environment
  - The agent improves its policy
- Planning:
  - A model of the environment is known
  - The agent performs computations with its model (without any external interaction)
  - The agent improves its policy a.k.a. deliberation, reasoning, introspection, pondering, thought, search

# Origins of RL



# RL Agent Taxonomy



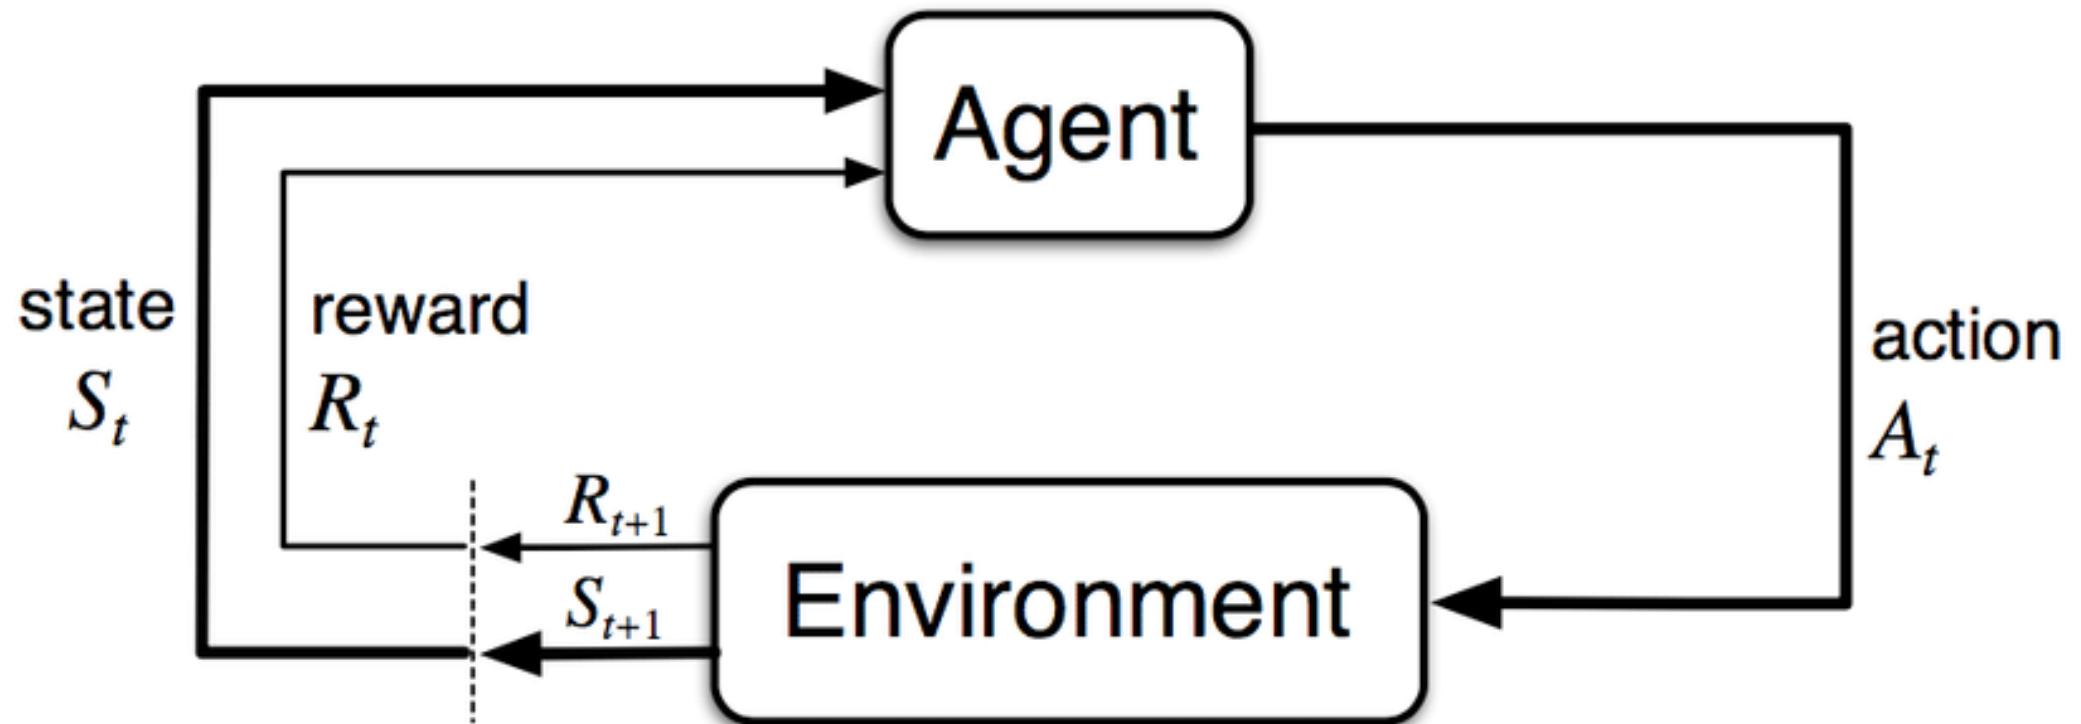
# Terminology

- Policy: agent's behavior function
  - Finding optimal policy known as the control problem
- Value function: how good is each state and/or action
  - Finding optimal value function is known as the prediction problem
- Model: agent's representation of the environment

# More Terminology on Types of RL

- Model free ← will build up to today
- Model based
- On Policy ← will build up to today
  - Learn directly from your experiences “on the job”
- Off policy
  - Learn from someone else’s behavior

# Markov Decision Process



# Formal MDP Definition

A Markov Decision Process is a tuple  $\langle S, A, T, R, \gamma \rangle$

- $S$  is a finite set of states
- $A$  is a finite set of actions
- $T$  is a state transition probability matrix,  
 $T_{ss'}^a = P [S_{t+1} = s' | S_t = s, A_t = a]$
- $R$  is a reward function,  $R_s^a = E [R_{t+1} | S_t = s, A_t = a]$
- $\gamma$  is a discount factor  $\gamma \in [0, 1]$ .

# Returns and Discounting

- The return  $G_t$  is the total discounted reward from time-step  $t$ .

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- The value of receiving reward  $R$  after  $k+1$  time-steps is  $\gamma^k R$
- $\gamma \approx 0$  is “myopic”,  $\gamma \approx 1$  is “far-sighted”
- Why discount?

- Mathematically convenient, avoids infinite returns
- Animal/human/investment banker’s behavior shows preference for immediate reward

# Formal Definition of Policy

- Distribution of action over states:  $\pi(a|s) = P [A_t = a | S_t = s]$
- Policy depends only on current state not history, this is the Markov property bit of MDP (how do people get around this for cases where history does matter)
- Theorem (abridged): There always exists an optimal policy for a given finite MDP. It follows the optimal value function.

# Formal Definition of Value Function

- State value: expected return starting from state  $s$ , and then following policy  $\pi$ 
  - $v_\pi(s) = E_\pi [G_t | S_t = s]$
- Action value: is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$ 
  - $q_\pi(s, a) = E_\pi [G_t | S_t = s, A_t = a]$

# Dynamic Programming

- Building up to RL first requires understanding Dynamic Programming
- Dynamic sequential or temporal component to the problem Programming optimizing a “program”, i.e. a policy
- A method for solving complex problems by breaking them down into subproblems
  - Solve the subproblems → Combine solutions to subproblems

# When to use DP

Dynamic Programming is a very general solution method for problems which have two properties:

- Optimal substructure:
  - Principle of optimality applies
  - Optimal solution can be decomposed into subproblems
- Overlapping subproblems:
  - Subproblems recur many times
  - Solutions can be cached and reused
- Markov decision processes satisfy both properties Bellman equation gives recursive decomposition Value function stores and reuses solutions

# Prediction vs Control

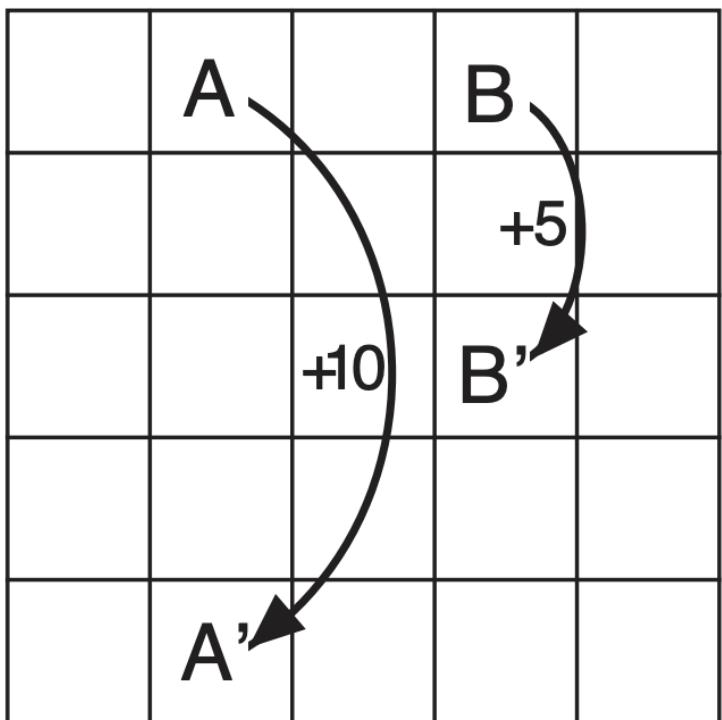
Two problems in RL

- Prediction is the problem of evaluating how good any given state is for getting rewards given a policy
- Control is the problem of selecting actions that give you a policy that maximizes reward

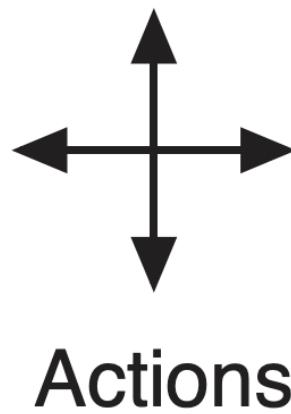
# Planning via DP

- Dynamic programming assumes full knowledge of the MDP
- It is used for planning in an MDP
- For prediction:
  - Input: MDP  $\langle S, A, T, R, \gamma \rangle$  and policy  $\pi$
  - Output: value function  $v_\pi$
- For control:
  - Input: MDP  $\langle S, A, T, R, \gamma \rangle$
  - Output: optimal value function  $v^*$  and: optimal policy  $\pi^*$

# Prediction Example



(a)



3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

(b)

# Bellman Expectation

- The state-value function can again be decomposed into immediate reward plus discounted value of successor state,

$$v_{\pi}(s) = E_{\pi} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]$$

- The action-value function can similarly be decomposed,

$$q_{\pi}(s, a) = E_{\pi} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

- No closed form solution (in general)

# Policy Evaluation

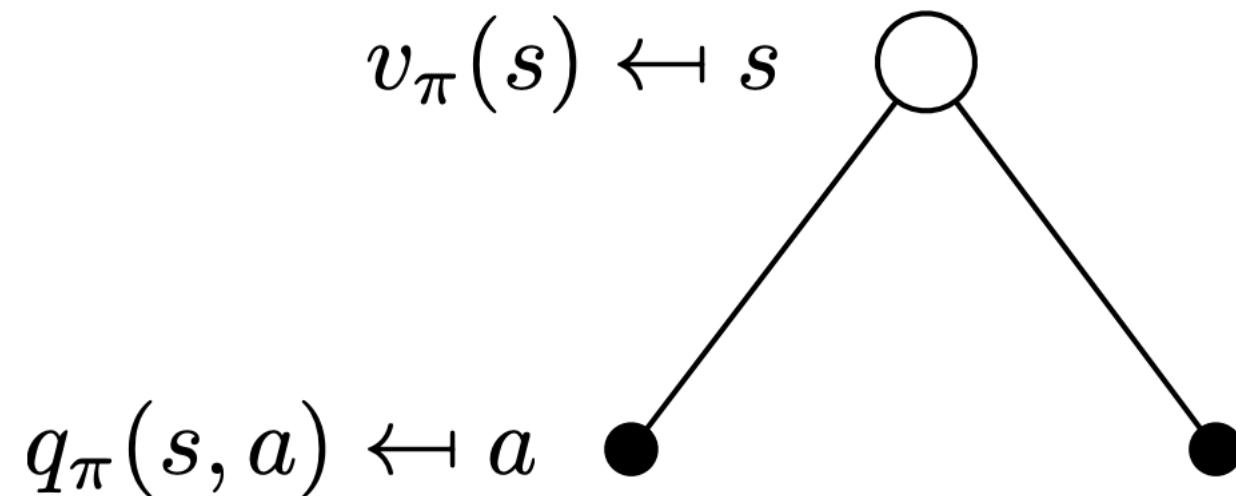
- Problem: evaluate a given policy  $\pi$
- Solution: iterative application of Bellman expectation backup

$$v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$$

- Using synchronous backups,
  - At each iteration  $k + 1$
  - For all states  $s \in S$  Update  $v_{k+1}(s)$  from  $v_k(s')$ , where  $s'$  is a successor state of  $s$

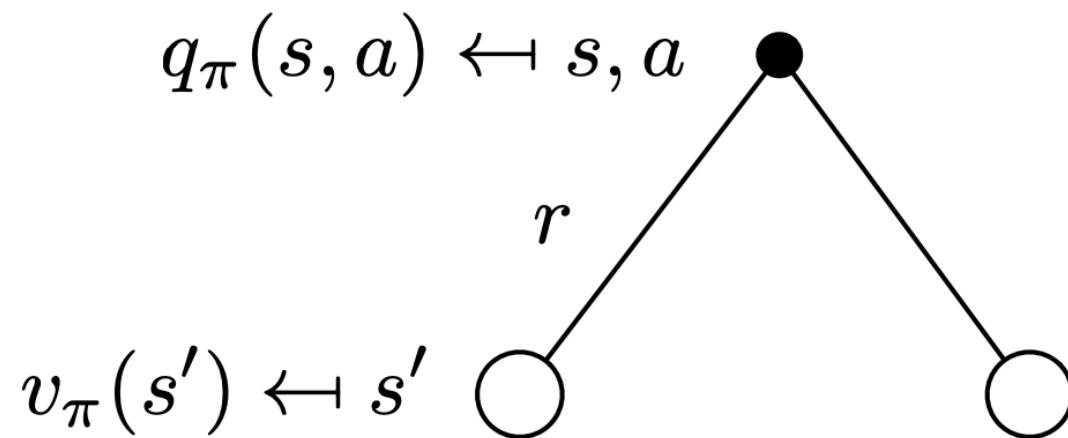
# Policy Evaluation

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s,a)$$



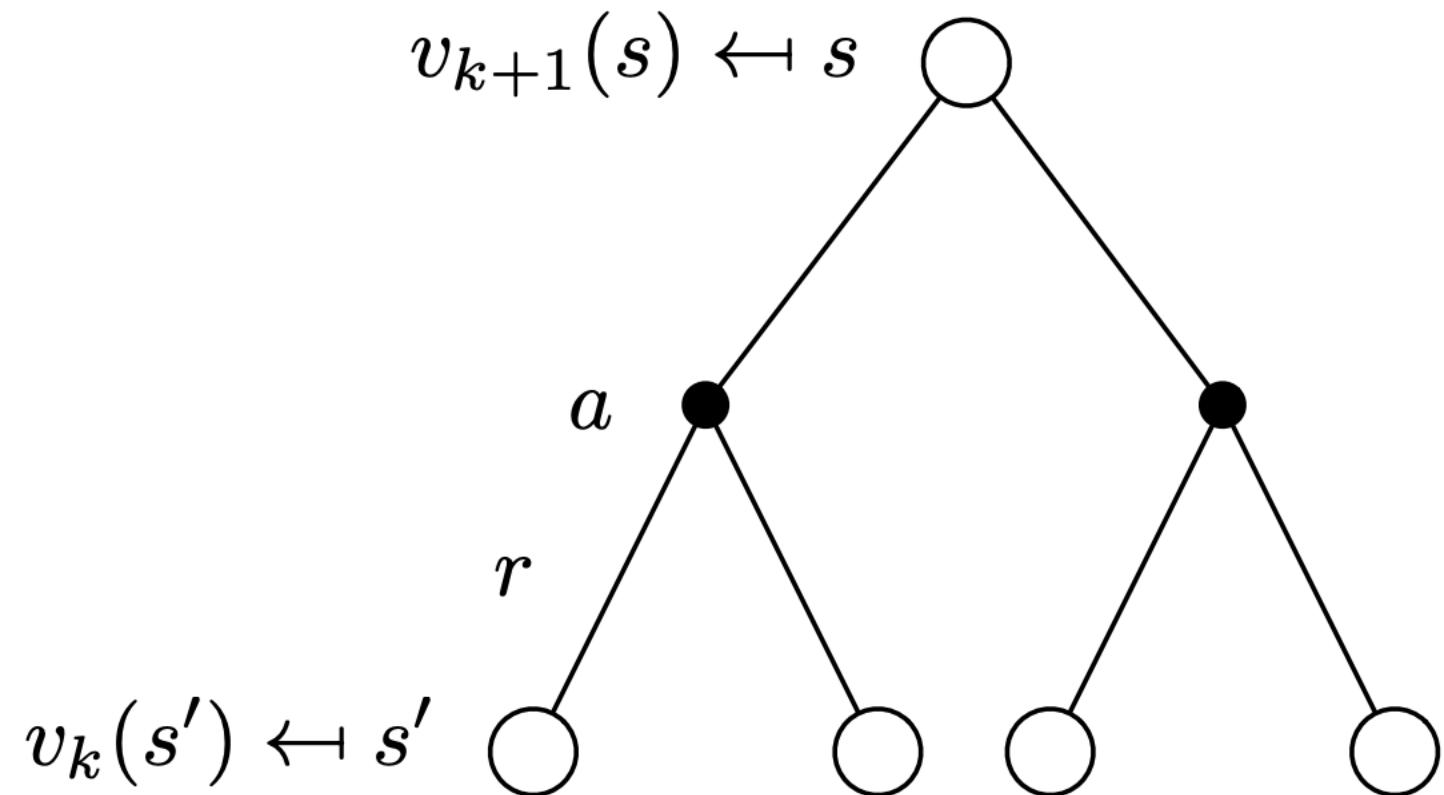
# Policy Evaluation

$$q_{\pi}(s, a) = R^a_s + \gamma \sum_{s' \in S} T^a_{ss'} v_k(s')$$

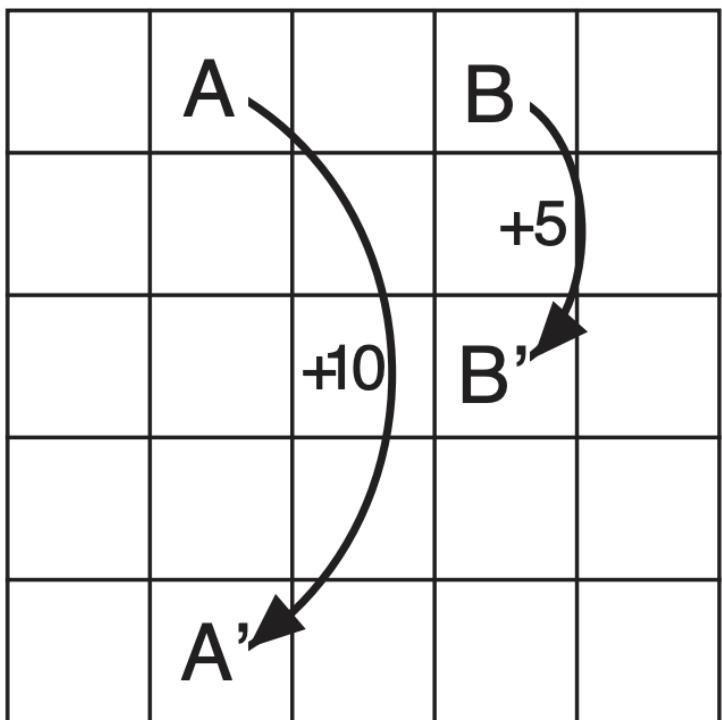


# Policy Evaluation

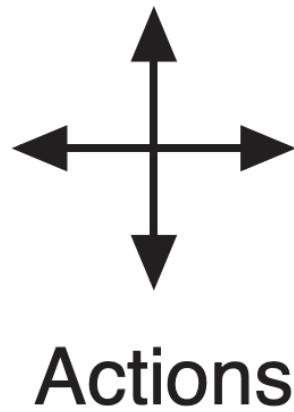
$$v_{k+1}(s) = \sum_{a \in A} \pi(a|s) (R^a_s + \gamma \sum_{s' \in S} T^a_{ss'} v_k(s'))$$



# Prediction Example



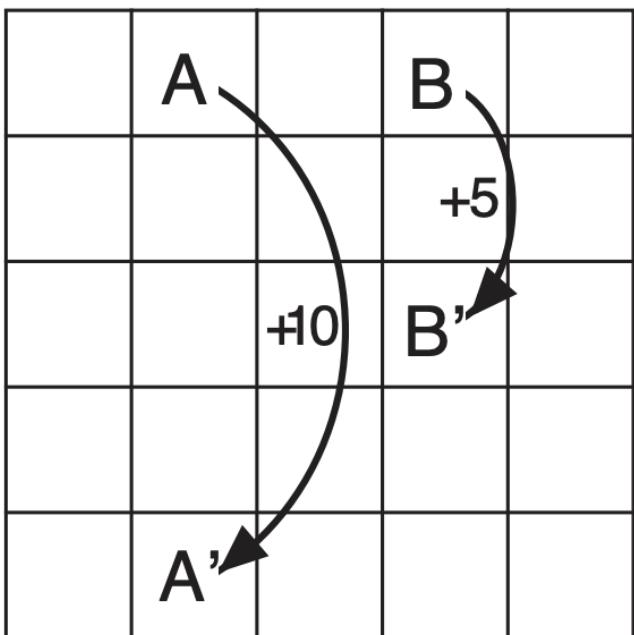
(a)



3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

(b)

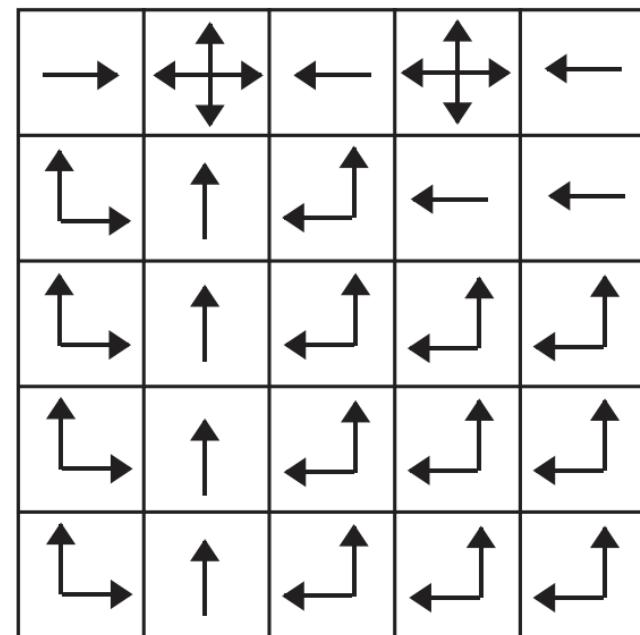
# Control Example



a) gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

b)  $v_*$



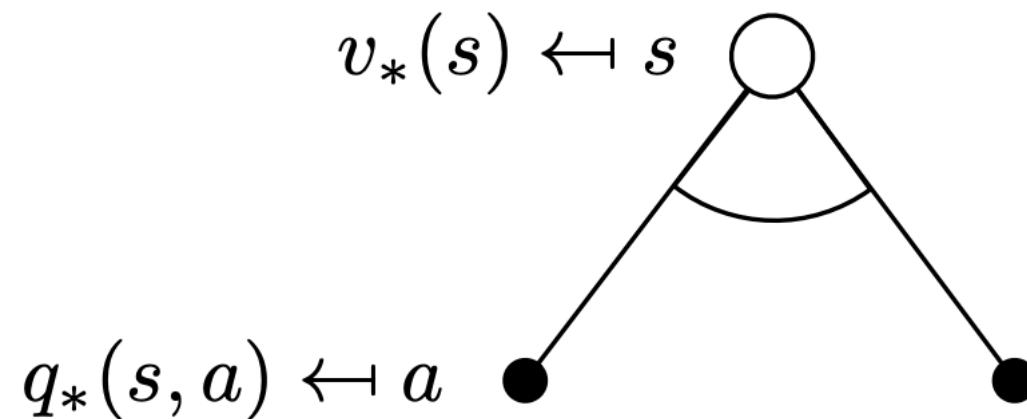
c)  $\pi_*$

# Bellman Optimality Equation

- Optimal state value:  $v^*(s) = \max_{\pi} v_{\pi}(s)$
- Optimal action value:  $q^*(s,a) = \max_{\pi} q_{\pi}(s,a)$
- Optimal policy:  $\pi^*(s) = \operatorname{argmax}_a q^*(s,a)$

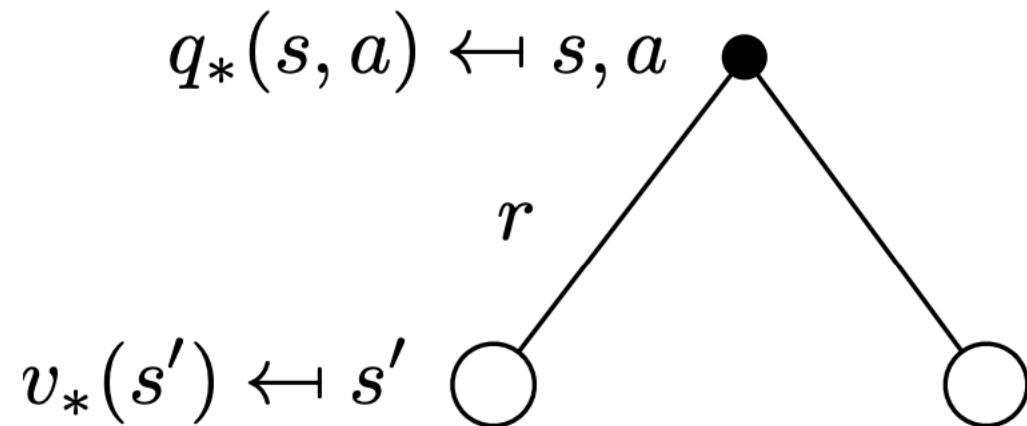
# Bellman Optimality Equation

$$v^*(s) = \max_a q^*(s', a')$$



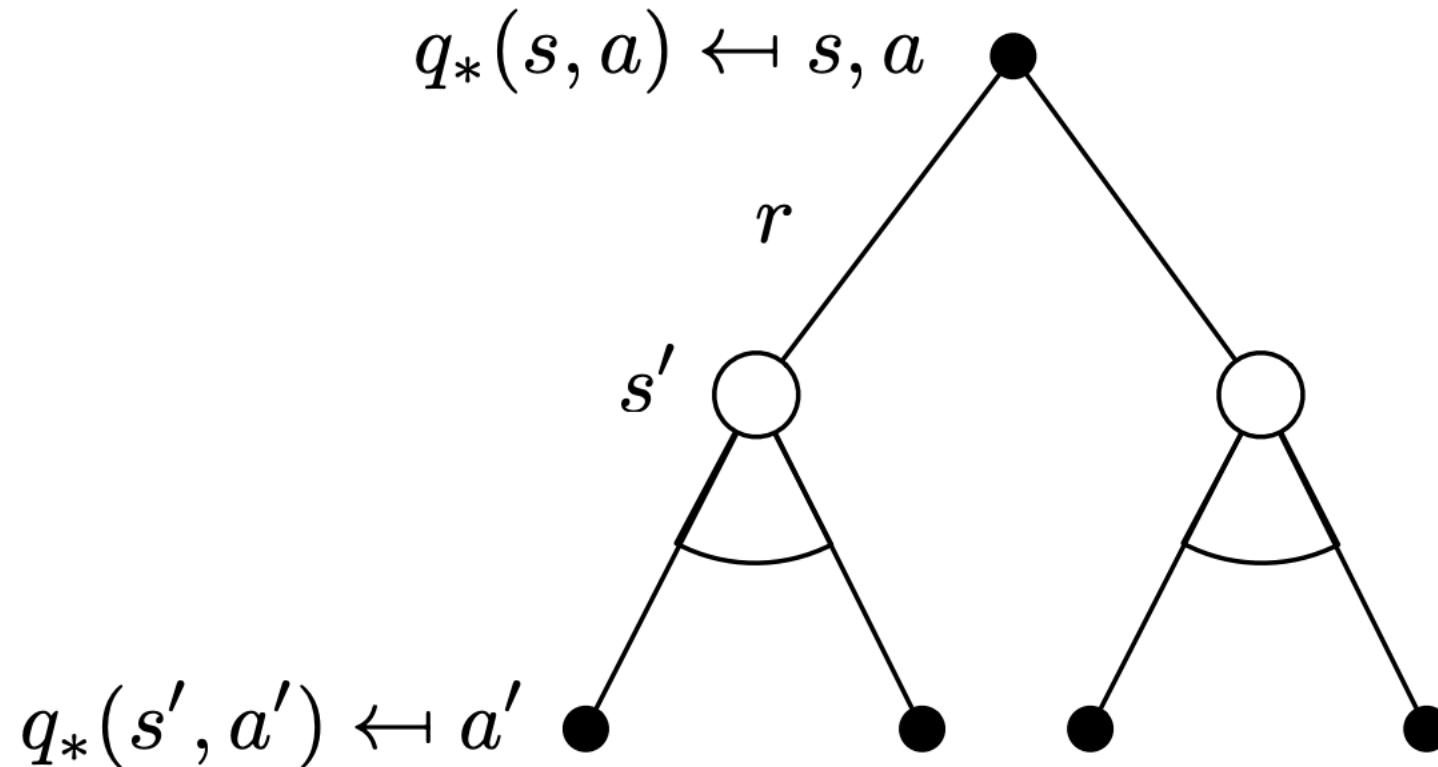
# Bellman Optimality Equation

$$q^*(s, a) = R_s^a + \gamma \sum_{s' \in S} T_{ss'}^a v^*(s')$$



# Bellman Optimality Equation

$$q^*(s, a) = R_s^a + \gamma \sum_{s' \in S} T_{ss'}^a \max_a q^*(s', a')$$



# Bellman Optimality Equation

- Optimal state value:  $v^*(s) = \max_{\pi} v_{\pi}(s)$
  - Optimal action value:  $q^*(s,a) = \max_{\pi} q_{\pi}(s,a)$
  - Optimal policy:  $\pi^*(s) = \operatorname{argmax}_a q^*(s,a)$
- 
- $q^*(s,a) = R_s^a + \gamma \sum_{s' \in S} T_{ss'}^a \max_a q^*(s',a')$
  - $v^*(s) = \max_{a \in A} q^*(s, a)$

# Policy Iteration

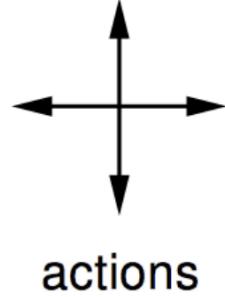
Given a policy  $\pi$

- Evaluate the policy  $v_\pi(s) = E [R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$
- Improve the policy by acting greedily with respect to  $v_\pi$   
 $\pi' = \text{greedy}(v_\pi)$
- Converting back and forth between prediction and control
- Start with random policy, eval it, improve value, improve policy

# Value Iteration

- Similar to Policy Iteration but start with random value function, recursively improve it
- Exercise to figure out equations if you start with random value instead of policy

# Put it together



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$   
on all transitions

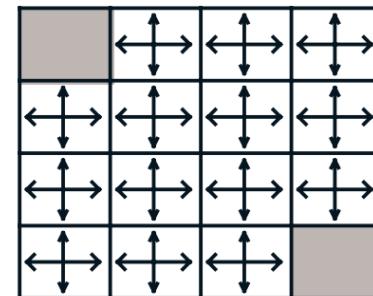
- Undiscounted episodic MDP ( $\gamma = 1$ )
- Nonterminal states 1, ..., 14
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is  $-1$  until the terminal state is reached
- Agent follows uniform random policy  $\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$

$v_k$  for the  
Random Policy

$k = 0$

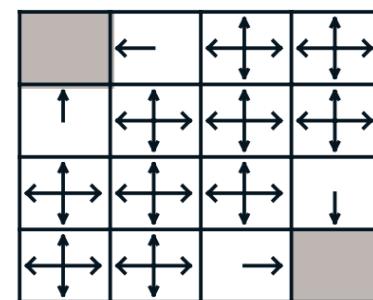
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

Greedy Policy  
w.r.t.  $v_k$



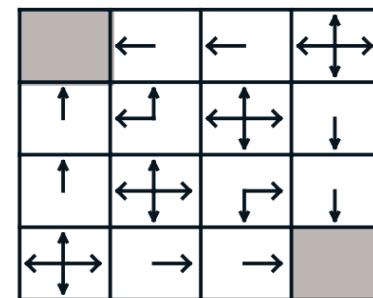
$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0



$k = 3$

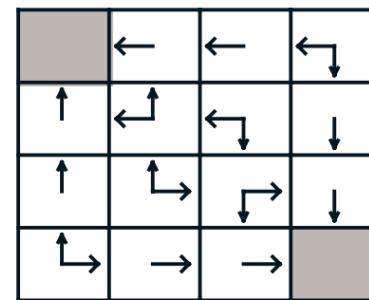
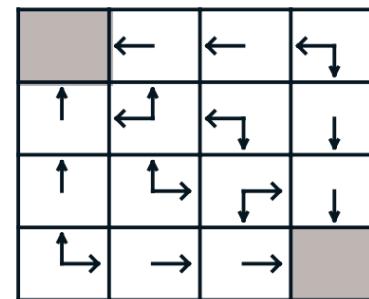
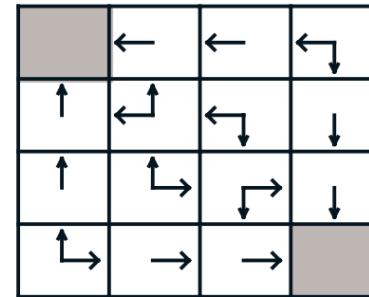
0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = \infty$

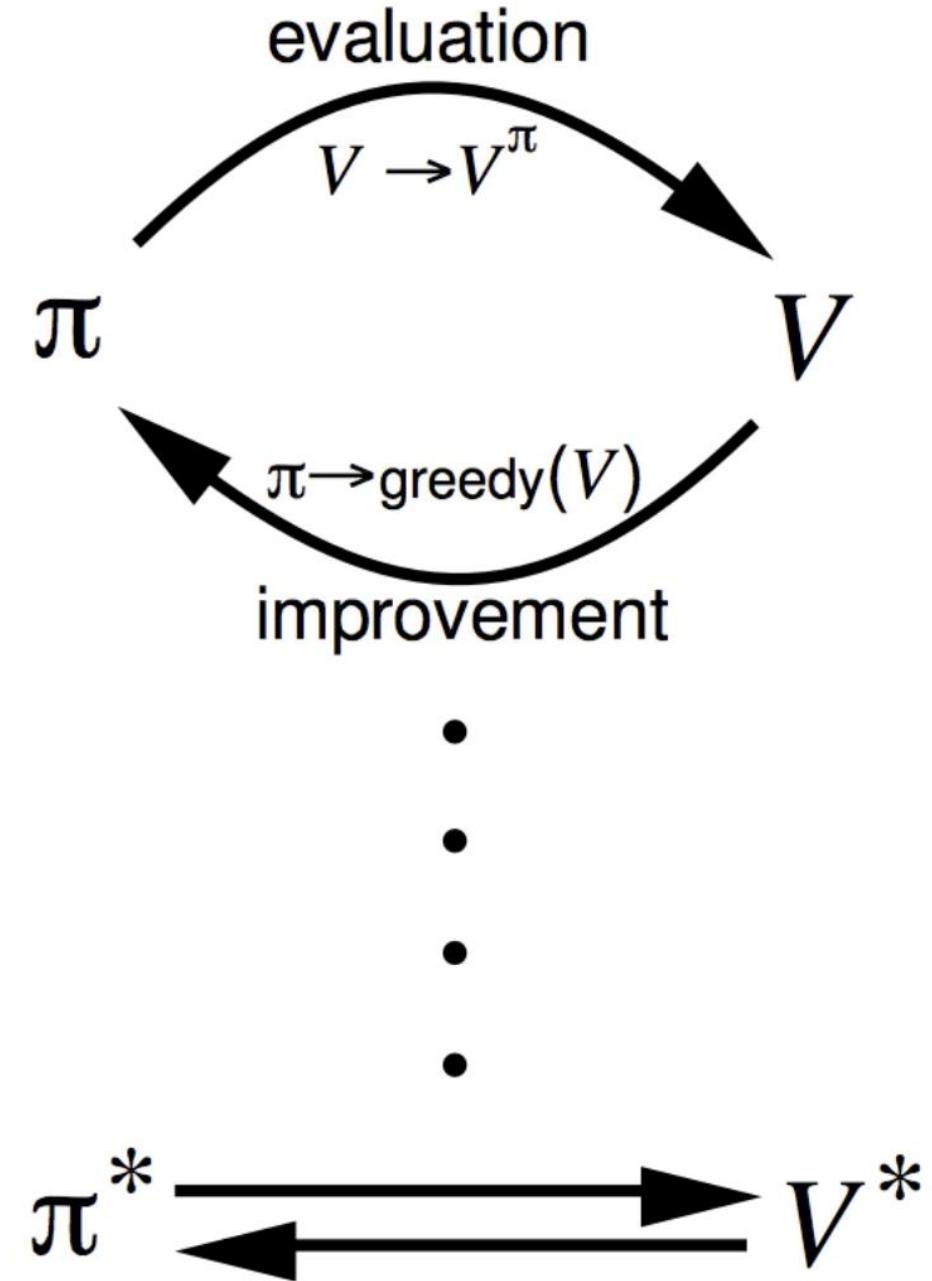
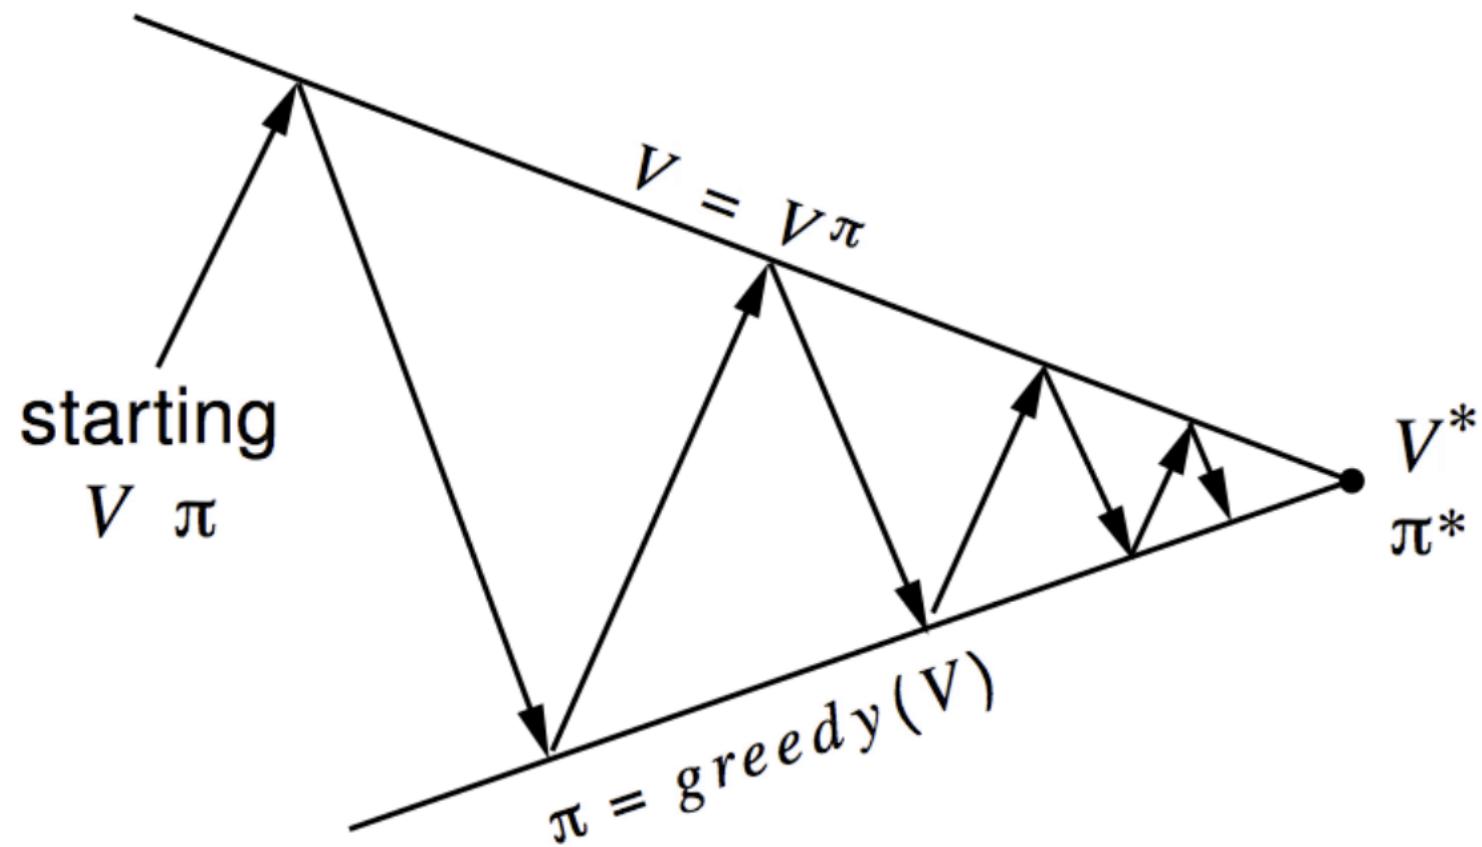
0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal  
policy

# Generalized Policy Iteration

- Both are iterative versions of this



# DP Limitations

- DP uses full-width backups
- For each backup Every successor state and action is considered
- Using knowledge of the MDP transitions and reward function DP is effective for medium-sized problems (millions of states)
- For large problems DP suffers Bellman's curse of dimensionality
- Number of states  $n = |S|$  grows exponentially with number of state variables Even one backup can be too expensive

# Model Free RL via Sample Backups

- Model Free RL: optimize value of unknown MDP
- Using sample rewards and sample transitions  $\langle S, A, R', S' \rangle$   
Instead of reward function  $R$  and transition dynamics  $T$
- Advantages: Model-free: no advance knowledge of MDP required  
Breaks the curse of dimensionality through sampling
- Cost of backup is constant, independent of  $n = |S|$

# Experience Based Learning

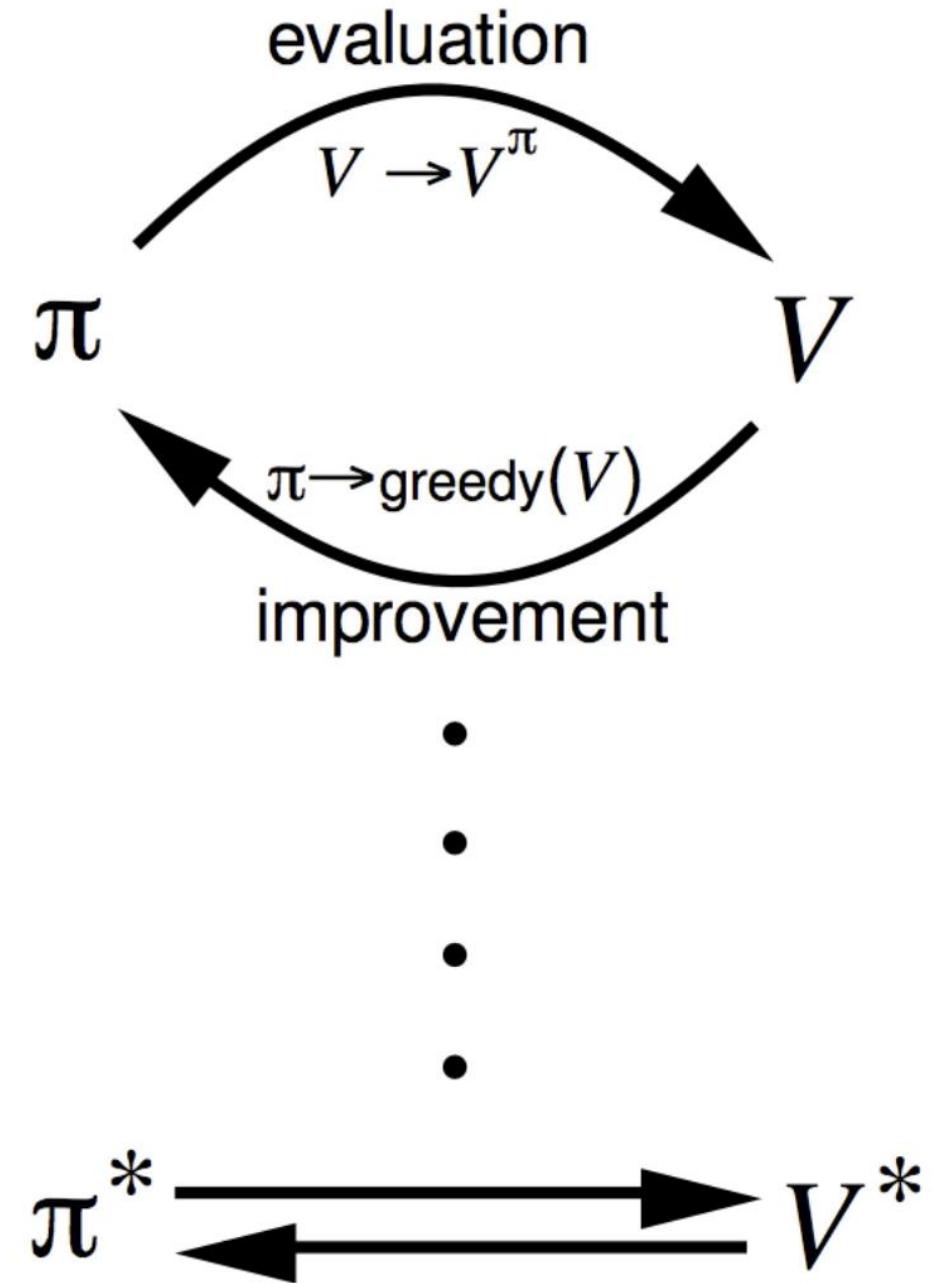
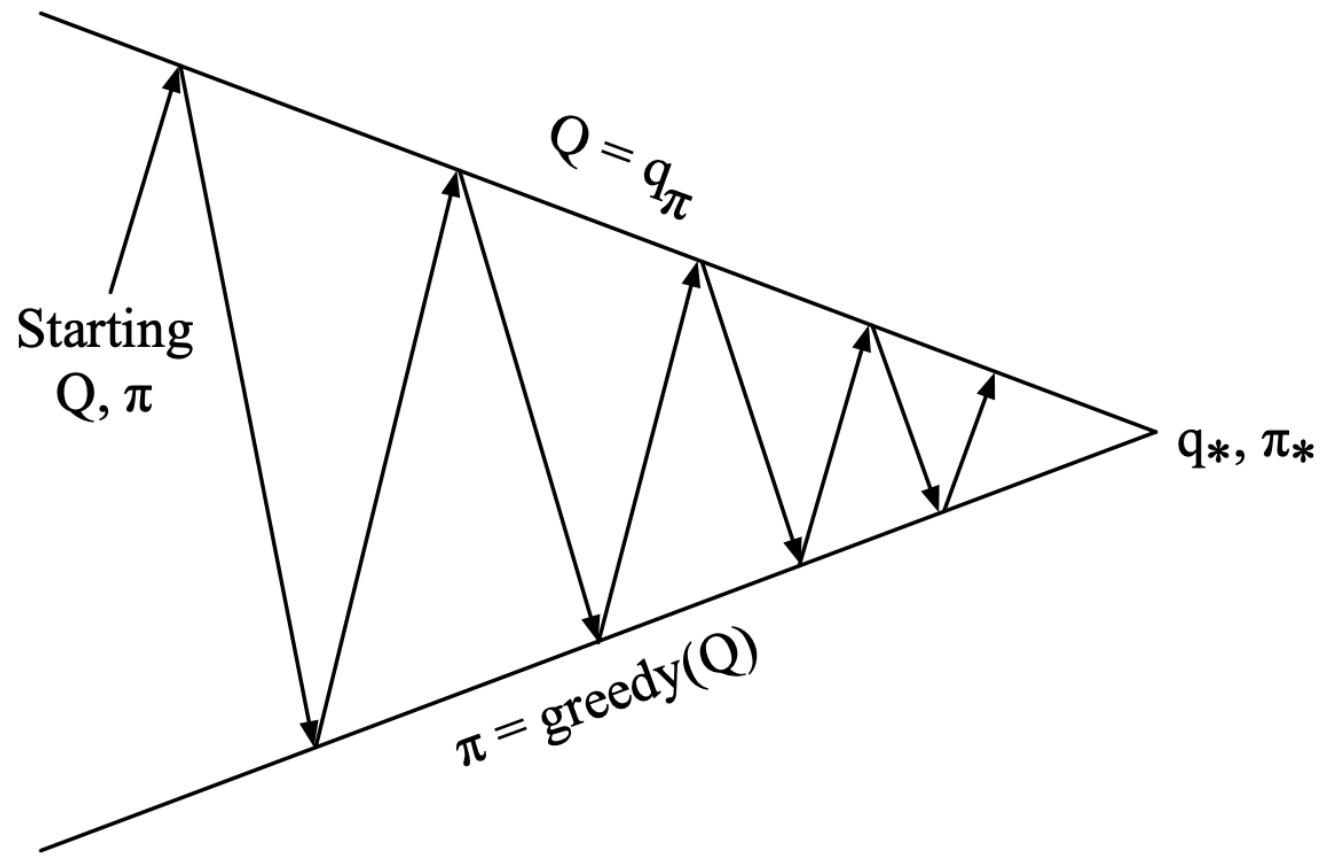
- Many real world problems are better suited to being solved by RL as opposed to DP based planning
- All the examples of agents we talked about first class
  - Robots in your home
  - Video games harder than tic tac toe
  - Language

# Monte Carlo Control

- Greedy policy improvement over  $V(s)$  requires model of MDP  
$$\pi'(s) = \operatorname{argmax}_{a \in A} R_s^a + \gamma \sum_{s' \in S} T_{ss'}^a V'(s')$$
- Greedy policy improvement over  $Q(s, a)$  is model-free  
$$\pi'(s) = \operatorname{argmax}_{a \in A} Q(s, a)$$
- Learn this  $Q$  by function approximation using the experiences you've gathered by Monte Carlo sampling

# Generalized Policy Iteration

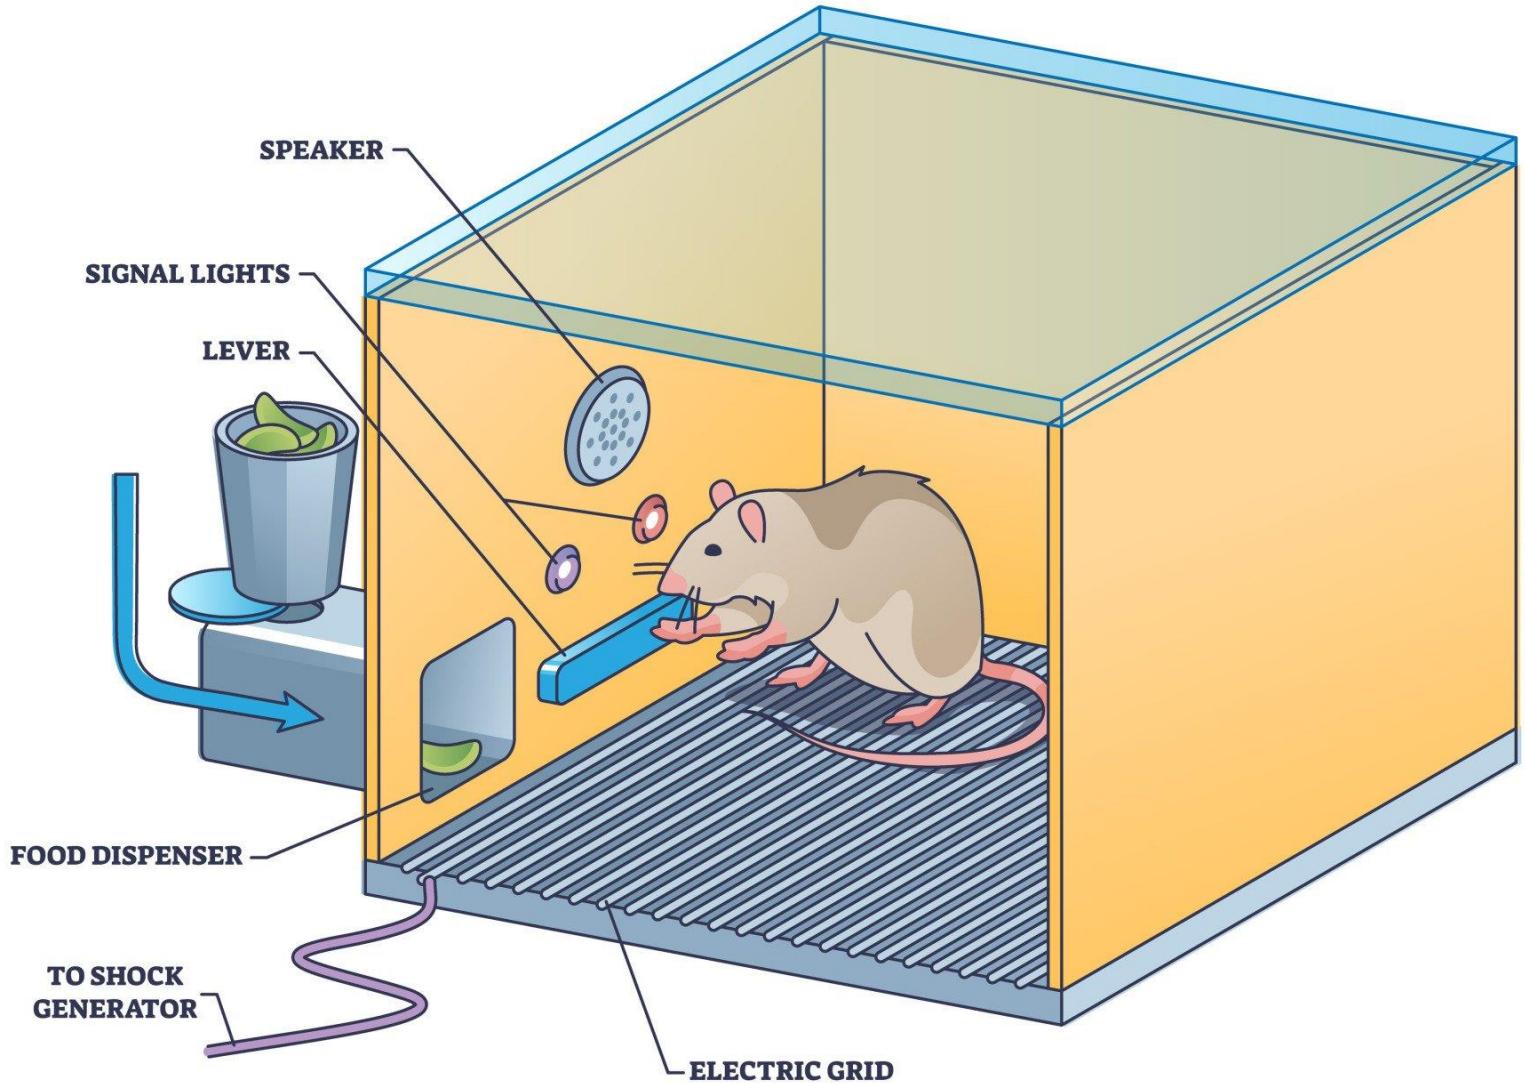
## Monte Carlo Evaluation



# Greedy Policy Improvement Limitations

- Greedy doesn't let you always explore all the actions you need

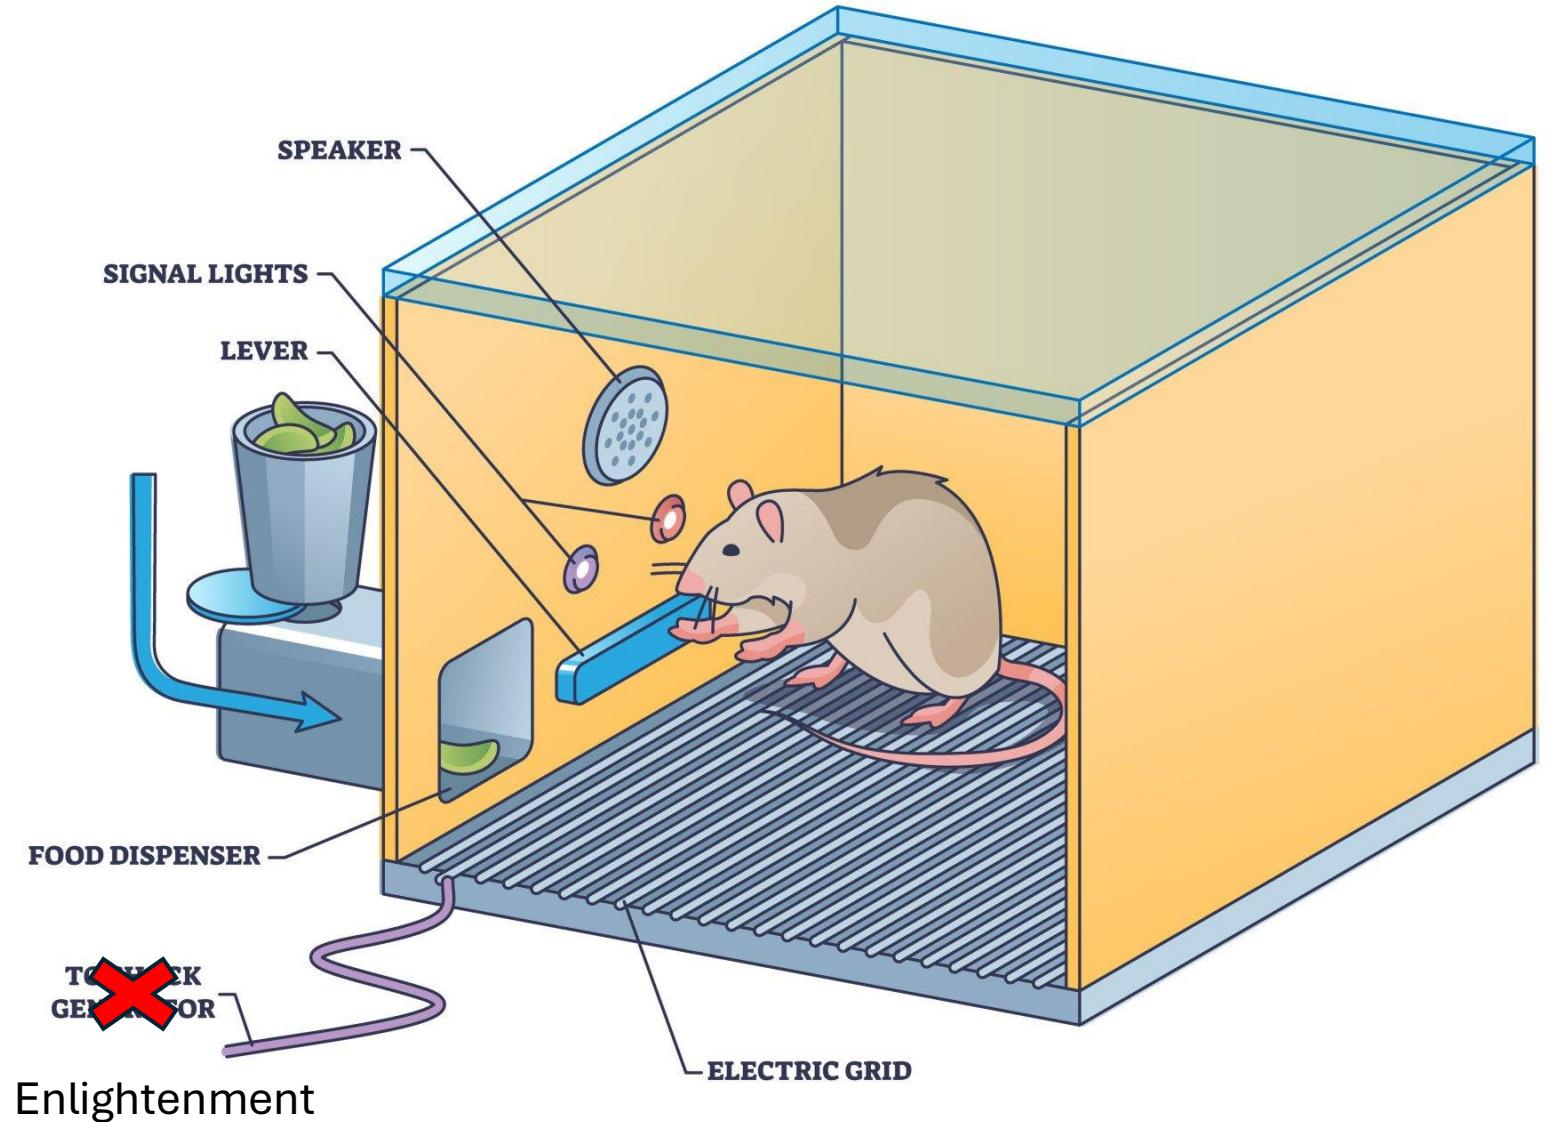
## SKINNER BOX



# Greedy Policy Improvement Limitations

- Greedy doesn't let you always explore all the actions you need

## SKINNER BOX



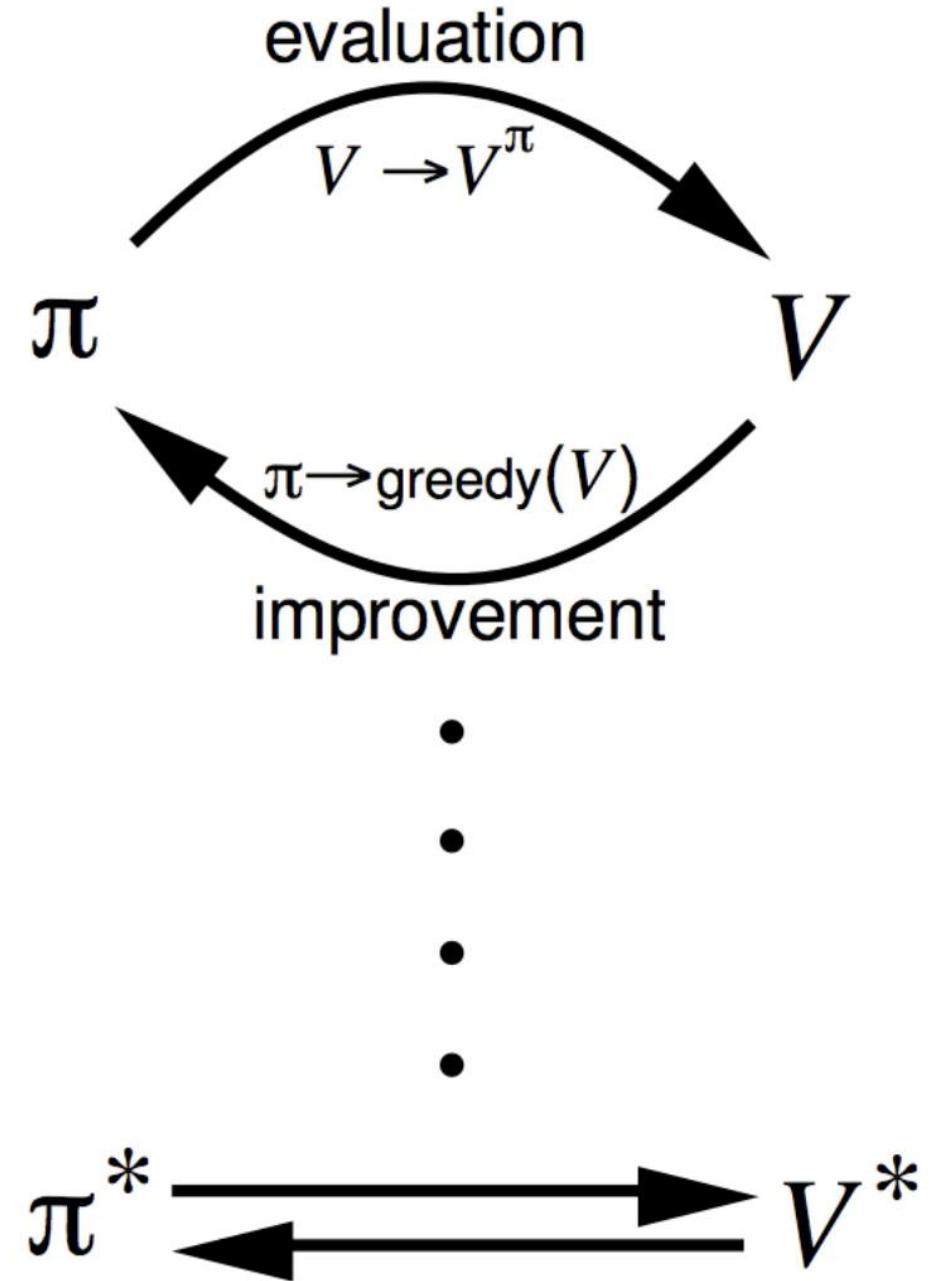
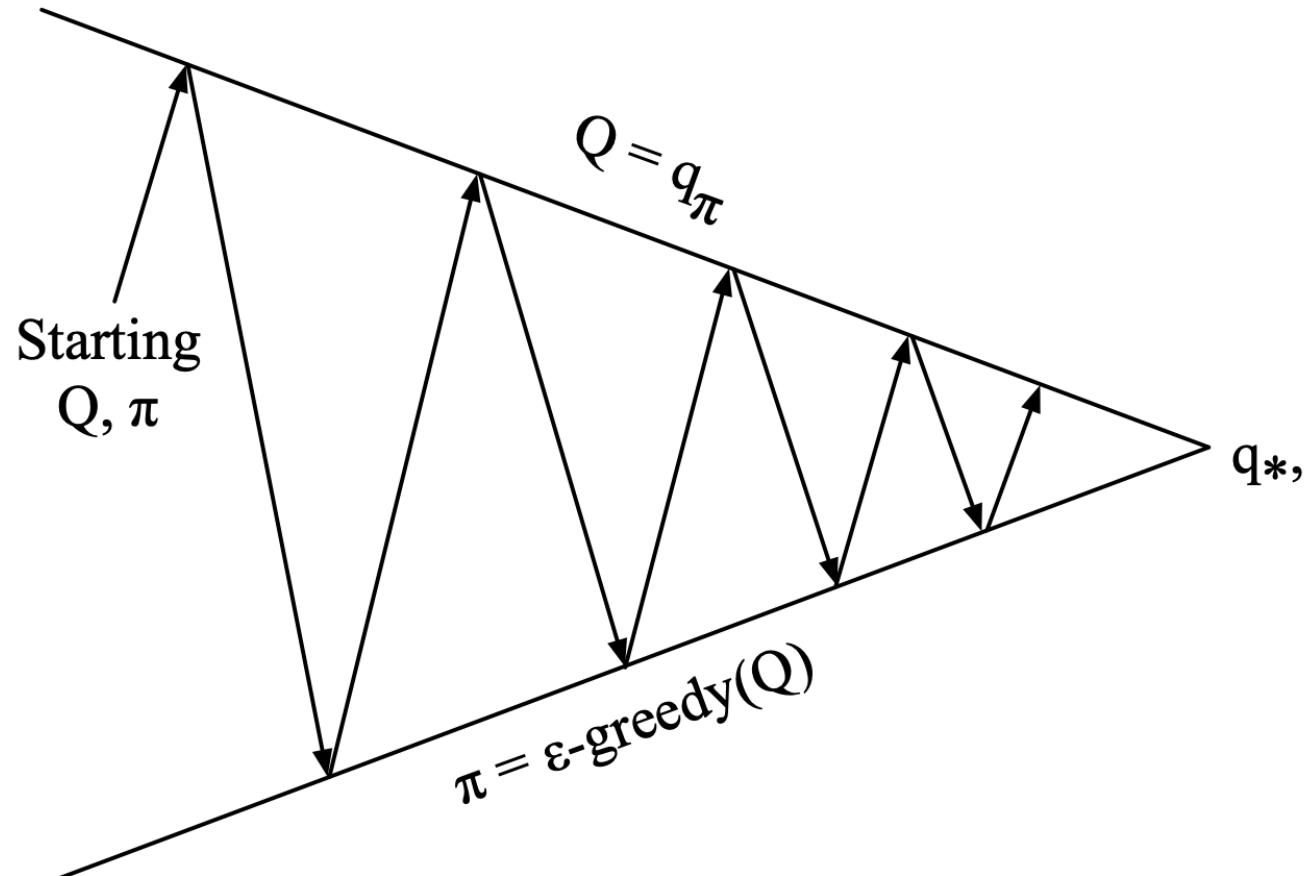
# $\epsilon$ -greedy exploration

- Simplest idea for ensuring continual exploration
- All  $m$  actions are tried with non-zero probability
- With probability  $1 - \epsilon$  choose the greedy action
- With probability  $\epsilon$  choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon & \text{if } a^* = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(s, a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

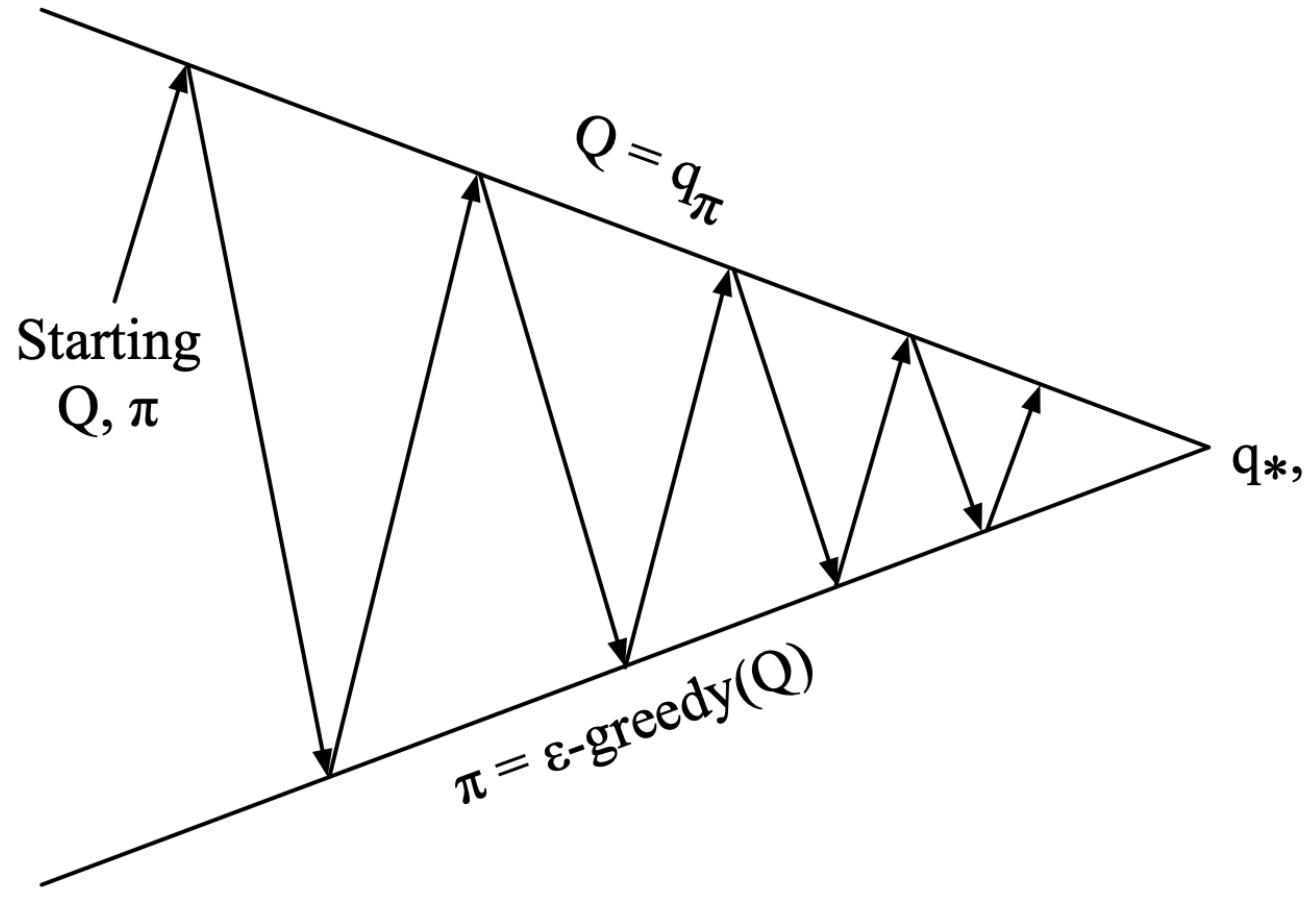
# Generalized Policy Iteration

## Monte Carlo Evaluation

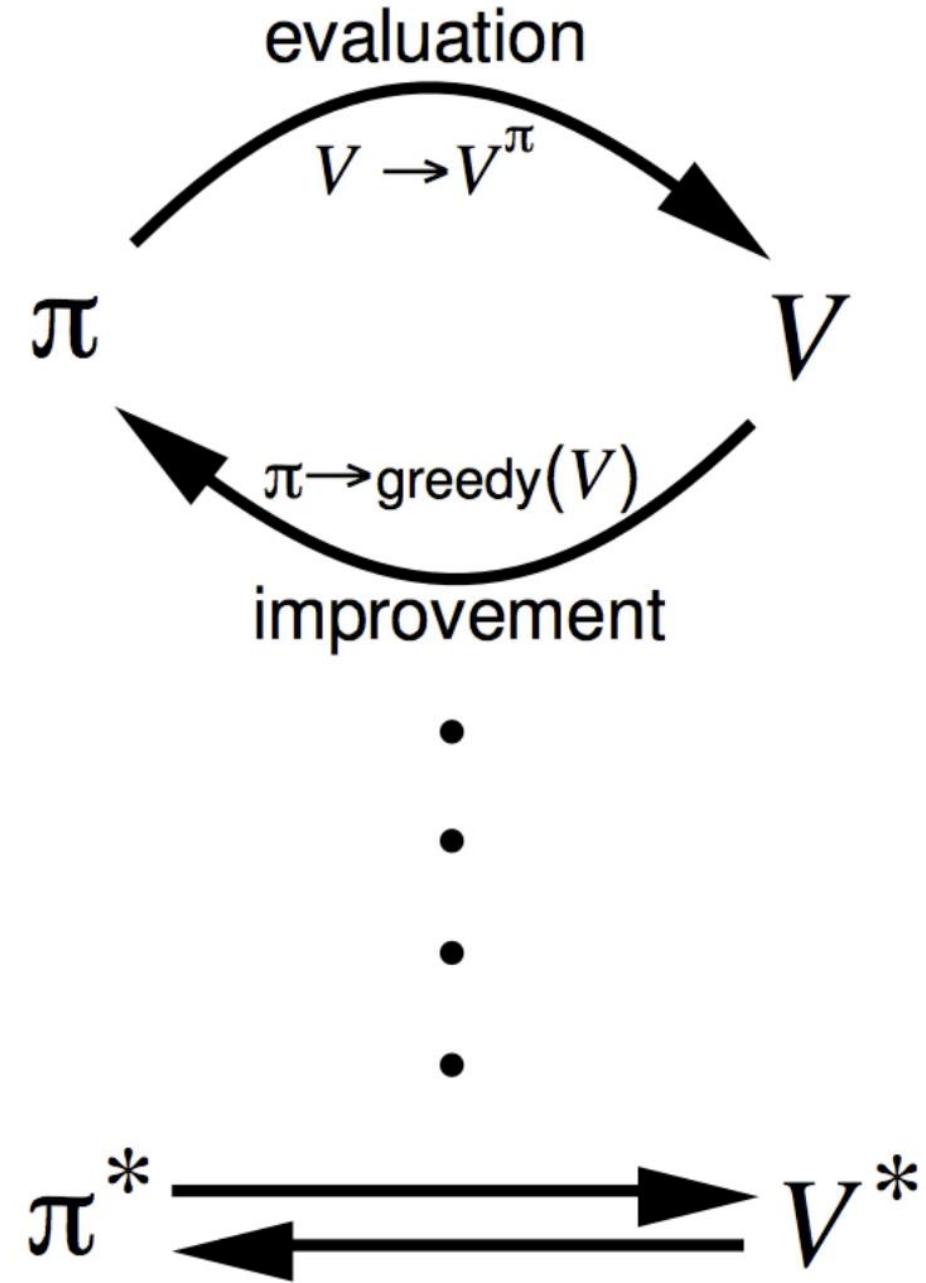


# Generalized Policy Iteration

## Monte Carlo Evaluation



You can't fully evaluate the entire state space each time



# Generalized Policy Iteration with Fn Approximation + Monte Carlo Eval

